π-Former: Succinct Proofs of Transformer Inference via Structured Lookup Attention

Hideaki Takahashi

April 2026

What Is a Transformer? (Intuition)

A transformer is a stack of L identical **blocks**. Each block takes a sequence of T token vectors $\mathbf{X} \in \mathbb{R}^{T \times d}$ and returns the same shape.

Inside one block, two operations:

1. **Attention**: each token looks at every other token and pulls in information from the relevant ones.
2. **Feed-forward (FFN)**: each token is independently passed through a small two-layer MLP.

Each is wrapped with: **LayerNorm** (rescale the row), a **residual** "+" (add the input back).

$$\mathbf{X}_{\text{mid}} = \mathbf{X}_{\text{in}} + \mathsf{Att}(\mathsf{LN}(\mathbf{X}_{\text{in}})) \qquad \mathbf{X}_{\text{out}} = \mathbf{X}_{\text{mid}} + \text{FFN}(\mathsf{LN}(\mathbf{X}_{\text{mid}}))$$

That's it. Repeat L times, project to vocabulary, predict next token.

What Is Attention? (Intuition)

Each token is given three vectors:

- **Query $\mathbf{q}$** – "what am I looking for?"
- **Key $\mathbf{k}$** – "what do I offer?"
- **Value $\mathbf{v}$** – "what info do I carry?"

all linear projections of the input row: $\mathbf{q} = \mathbf{x}\mathbf{W}_Q$, etc.

For token t, attention computes a *weighted average of values*, where the weight is how well its query matches each key:

$$\mathbf{out}_t = \sum_s \underbrace{\mathbf{softmax}_s\big(\langle \mathbf{q}_t, \mathbf{k}_s \rangle / \sqrt{d}\big)}_{\text{attention weight}} \mathbf{v}_s .$$

What this gives the model. Tokens can "ask" for information from other positions and aggregate it – no fixed window, no recurrence.

What this costs the SNARK. Softmax means $\exp(\cdot)$ and row-wise division – both *very* expensive in a finite field. This is what we will replace.

The Other Two: GeLU and LayerNorm

GeLU. A smooth nonlinearity used inside the FFN, applied element-wise:

$$\mathrm{GeLU}(x) \;=\; x \cdot \Phi(x), \qquad \Phi = \text{Gaussian CDF}.$$

Like a soft ReLU. Lets the FFN learn nonlinear features.
SNARK pain: transcendental (CDF / exp).

LayerNorm. Stabilises training by rescaling each row to zero-mean / unit-variance, then re-scaling with learned γ, β:

$$\mathsf{LN}(\mathbf{x}) \;=\; \gamma \odot \frac{\mathbf{x} - \mu(\mathbf{x})}{\sqrt{\sigma^2(\mathbf{x}) + \epsilon}} + \beta.$$

Without it, deep stacks blow up.
SNARK pain: a square root and a division per row.

So a transformer block has *four* hostile pieces

softmax, GeLU, LayerNorm, dense $\mathbf{W}$ – each adds either non-arithmetic operations or a field multiplication per weight.

Motivation, Challenge, and Our Angle

Goal. Prove $\mathbf{y} = \mathcal{M}_\theta(\mathbf{x})$ with a short proof and cheap verification.

Catch. A SNARK only natively expresses additions and multiplications in $\mathbb{F}$. State of the art (zkLLM, zkGPT) approximates the four hostile ops; prover time is still minutes per token at GPT scale.

Our angle: co-design. The narrative is driven by the dense projections, which dominate the prover's cost:

Standard op	$\Rightarrow$	**π-Former op**
dense $\mathbf{W}$ (FP)	$\Rightarrow$	**ternary** $\mathbf{W} \in \{-1, 0, +1\} \cdot \alpha$ *(no field mults)*
softmax attention	$\Rightarrow$	**LLA**: linear attention with learnable $\boldsymbol{\phi}$
GeLU	$\Rightarrow$	same learnable additive sub-table lookup
LayerNorm	$\Rightarrow$	constraint-fused LN *(zkGPT-style)*

Toolbox (One-Line Reminders)

Sumcheck. Reduces $H = \sum_{x \in \{0,1\}^n} f(x)$ to a single evaluation $f(r)$ at random r. Verifier cost $O(n\delta)$, soundness error $\leq n\delta/|\mathbb{F}|$.

Hyrax PCS. Transparent (no trusted setup) polynomial commitment based on Pedersen / DL on BN254 G1. $O(\sqrt{N})$ commitment size, $O(\sqrt{N})$ open. Importantly: **additively homomorphic**.

Lasso. A lookup argument that proves $v_j = T[\mathrm{idx}_j]$ for many queries against a table T, *cheap when T decomposes into small additive sub-tables*.

Strategy

Build the model so that: attention = sumcheck + Lasso, activation = Lasso, LayerNorm = sumchecks + range checks, projection = sumcheck with no field mults.

The π-Former Block

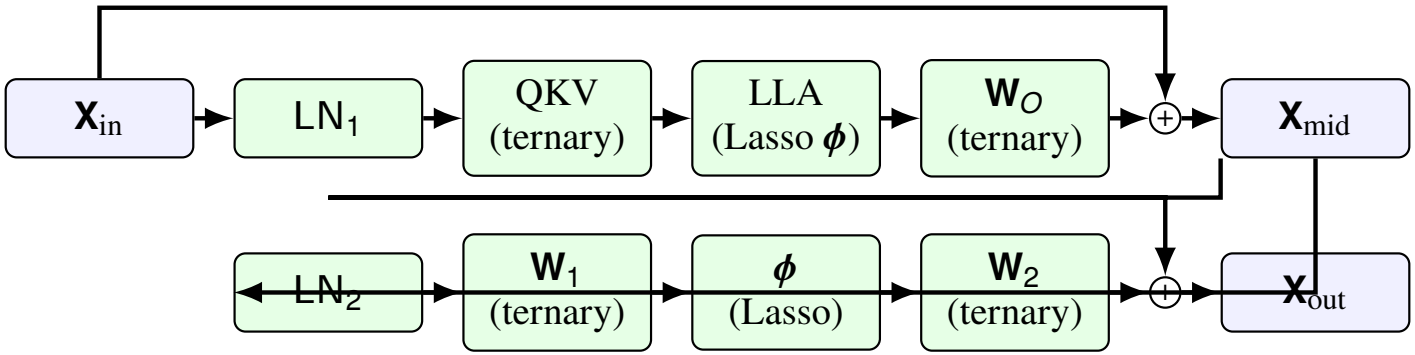

Color code.

- Green = an op the SNARK proves.
- Ternary boxes have *no* field multiplications inside.
- LLA / $\boldsymbol{\phi}$ are proved via Lasso lookups.
- "+" (residual) is checked by **homomorphic addition of Hyrax commitments** – no proof at all.

Novel Idea 1: Learnable Lookup Attention (LLA)

Linear attention replaces $\mathrm{softmax}(\mathbf{Q}\mathbf{K}^\top/\sqrt{d})$ by a positive-definite kernel $\kappa(\mathbf{q},\mathbf{k}) = \langle \boldsymbol{\phi}(\mathbf{q}), \boldsymbol{\phi}(\mathbf{k})\rangle$:

$$\mathrm{LLA}(\mathbf{Q},\mathbf{K},\mathbf{V})_t \;=\; \boldsymbol{\phi}(\mathbf{q}_t)^\top \underbrace{\sum_s \boldsymbol{\phi}(\mathbf{k}_s)\mathbf{v}_s^\top}_{=\mathbf{M}\ \text{(context, computed once)}} .$$

Cost: $O(n\,d^2)$ instead of $O(n^2 d)$, no exp, no row division.

Where's $\boldsymbol{\phi}$? We make it learnable and SNARK-friendly:

$$\boldsymbol{\phi}(x) \;=\; \sum_{i=0}^{c-1} T_i\big[\chi_i(x_{\mathrm{int}})\big],$$

i.e. split a b-bit input into c chunks, look each up in a small table T_i of size $2^{b/c}$, sum.

Why this matters.

- For b=16, c=2: two 256-entry tables instead of one 65,536-entry table.
- This is *exactly* the structured-decomposability condition Lasso wants.
- Init so $\sum_i T_i \approx$ GeLU, then train end-to-end with the straight-through estimator.

Novel Idea 2: Ternary Projections

All projection matrices ($\mathbf{W}_Q, \mathbf{W}_K, \mathbf{W}_V, \mathbf{W}_O, \mathbf{W}_1, \mathbf{W}_2$, LM head) are constrained:

$$\mathbf{W} \in \{-1, 0, +1\}^{m \times n}, \qquad \mathbf{y} = \alpha \cdot (\mathbf{x}\mathbf{W}) + \mathbf{b}$$

with one learnable scalar α per layer (BitNet b1.58 style).

What changes inside the SNARK? Consider the standard sumcheck for a contraction $\mathbf{Y} = \mathbf{X}\mathbf{W}$:

$$\widetilde{\mathbf{Y}}(r_t, r_d) - \mathbf{b}(r_d) \;=\; \alpha \sum_k \widetilde{\mathbf{X}}(r_t, k)\, \widetilde{\mathbf{W}}(k, r_d).$$

Because $\widetilde{\mathbf{W}}(k, r_d) \in \{-1, 0, +1\}$ after evaluation, each prover round is:

"+= / −= / skip" – no field multiplications.

Bonus

The verifying key never holds raw weights – only Hyrax *commitments* to the ternary arrays.

Re-using zkGPT-style Constraint-Fused LayerNorm

Not our novel idea. Constraint fusion across a transformer block was introduced in **zkGPT**; we adopt and re-implement it here so that the LN step matches our sumcheck + Lasso pipeline.

What it does. Rewrite the LN identity as integer relations, prove them at *one* random row index r_t in four steps:

1. **Mean**: degree-2 sumcheck.
2. **Variance**: cubic sumcheck on the centered residual.
3. **Sigma range**: prove $\sigma = \lfloor \sqrt{\mathrm{var}}/d \rfloor$ via two range checks (no $\sqrt{}$ in the circuit).
4. **Y fusion**: one multi-batched cubic sumcheck for the $\gamma\mathbf{X}$ and $\sigma\mathbf{Y}$ legs together.

Result. No field divisions, no square roots. Verifier reads γ, β from the VK; everything else is a Hyrax opening.

Our novel pieces are LLA + ternary projections

LayerNorm is the SNARK-friendly piece we *re-use* from prior work.

Optimization: GKR and Batching

Two orthogonal collapses make the proof short:

(A) Cross-block batching *(over L blocks)*. A transformer is the *same protocol run L times*. After absorbing per-block claims $\{H_\ell\}$, draw $\eta \leftarrow \mathsf{FS}$ and prove

$$H^\star = \sum_{\ell=1}^{L} \eta^{\ell-1} H_\ell$$

with one inner sumcheck. $\Rightarrow$ 6 model-level sumchecks (QKV, $\mathbf{W}_O$, attn-out, attn-ctx, FFN-Y, FFN-M) instead of $6L$. Soundness cost: only $(L-1)/|\mathbb{F}|$.

(B) GKR-style combine *(per polynomial)*. The same committed witness f (e.g. $\mathbf{X}_{\mathrm{n1}}$) is evaluated at k different random points $z_1, \ldots, z_k$ across protocols. Naively that is k Hyrax openings on f. Instead, with FS weights r_i, run *one* sumcheck for

$$\sum_x G(x) \cdot f(x), \qquad G(x) = \sum_{i=1}^{k} r_i \, \widetilde{\mathrm{eq}}(z_i, x),$$

ending in **one** Hyrax opening of f at r_{final}. The verifier computes $G(r_{\mathrm{final}})$ locally in $O(k \cdot n)$ field ops; no PCS opening for G is needed.

End-to-End Performance

Single-thread x86_64 CPU. Times are pure compute; IO excluded.

Model	d	L	d_{ff}	T	**Params**	**Prove**	**Verify**	**Proof**
tiny	8	1	16	8	1.1k	0.22 s	0.03 s	108 KB
small	64	2	256	16	133k	1.10 s	0.15 s	1.0 MB
medium	128	4	512	32	924k	**4.24 s**	**0.28 s**	5.5 MB

Headline. 4-layer, ~924k-parameter transformer: **4.2 s prove, 280 ms verify, 5.5 MB proof** on a single CPU core.

Scaling.

- Prove time ~ linear in $L \cdot T \cdot d$ (per-block arithmetic survives batching).
- Verify time *sub-linear*: each cross-block sumcheck is a $\log Td$-round transcript, accumulators finalise in 1 MSM pair each.

Comparison, Limitations, Roadmap

System	Attention	Activation
zkLLM	softmax (`tlookup`)	structured lookup
zkGPT	softmax (approx., Plonky2)	lookup table
π-Former	**linear (LLA, exact)**	**learnable** additive lookup

Position. π-Former attention is not a softmax approximation: it is a different (kernelised) attention trained from scratch, so the SNARK proves the model *exactly*. $\boldsymbol{\phi}$ is learnable.

Currently in the SNARK. Single-head; un-normalised LLA numerator; no causal mask; *not* zero-knowledge.

Engineering, not architecture.

- Multi-head: parallel batches of the same sumcheck.
- Causal mask: an indicator polynomial inside the sumcheck.
- ZK: standard blinding layer over committed witnesses.
- GPU MSM + parallel sumcheck: every dominant inner loop is data-parallel.
- Streaming generation: Nova-style folding with $\mathbf{M} = \boldsymbol{\phi}(\mathbf{K})^\top \mathbf{V}$ as the IVC accumulator $\Rightarrow$ constant amortised cost per token.

Take-aways

1. **Co-design beats approximation.** Replace softmax / GeLU / LN / dense $\mathbf{W}$ with twins that are *natively* sumcheck + Lasso.
2. **Our novel pieces**: LLA (linear, lookup-based attention) and ternary projections (no in-circuit field mults). LN re-uses zkGPT-style fusion.
3. **Cross-block batching** folds L per-block proofs into 6 model-level sumchecks, 1 global Lasso, 1 global range, 10 deferred openings.
4. **Result.** 4-layer, 924k-parameter transformer: **4.2 s prove / 280 ms verify / 5.5 MB proof**, single CPU core, transparent setup.

Thank you.

ht2673@columbia.edu